

Docker Dojo

Doug Davis dug@us.ibm.com

Morgan Bauer mbauer@us.ibm.com



Agenda

- Introduction to Containers & Docker
- DevOps: Configuration Management
- DevOps: Building Value
- DevOps: Scaling and Reliability



Preface

- Our goal: Learn about Docker
- But, as we learn about Docker, ask yourself:
 - Does Docker make the DevOps end-to-end experience better?
 - How can a customer's existing workflow leverage this technology?
 - Can do we this with VMs?
 - Ask lots of questions!



Introduction to Containers and Docker



What is Docker?

- At its core, Docker is tooling to manage containers
 - Docker is not a technology, it's a tool or platform
 - Simplified existing technology to enable it for the masses
- But, let's first discuss containers...




What are Containers?

- A group of processes run in isolation
 - Similar to VMs but managed at the process level
 - All processes MUST be able to run on the shared kernel
- Each container has its own set of "namespaces" (isolated view)
 - **PID** - process IDs
 - **USER** - user and group IDs
 - **UTS** - hostname and domain name
 - **NS** - mount points
 - **NET** - Network devices, stacks, ports
 - **IPC** - inter-process communications, message queues
 - **cgroups** - controls limits and monitoring of resources
- Docker gives it its own root filesystem



VM vs Container


Virtual Machine




Hypervisor

Hardware

Each VM has its own OS



Container



Base OS/Kernel

Hardware

Containers share the same base Kernel

App, bins/libs/OS must all be runnable on the shared kernel

If OS files aren't needed they can be excluded.

< VM ?

Why Containers?

- Fast startup time - only takes milliseconds to:
 - Create a new directory
 - Lay-down the container's filesystem
 - Setup the networks, mounts, ...
 - Start the process
- Better resource utilization
 - Can fit far more containers than VMs into a host



What is Docker again?

- **Tooling** to manage containers
 - Containers are not new
 - Docker just made them easy to use
- Docker creates and manages the lifecycle of containers
 - Setup filesystem
 - CRUD container
 - Setup networks
 - Setup volumes / mounts
 - Create: start new process telling OS to run it in isolation



Our First Container

```
$ docker run ubuntu echo Hello World  
Hello World
```

- What happened?
 - Docker created a directory with a "ubuntu" filesystem (image)
 - Docker created a new set of namespaces
 - Ran a new process: `echo Hello World`
 - Using those namespaces to isolate it from other processes
 - Using that new directory as the "root" of the filesystem (`chroot`)
 - That's it!
 - Notice as a user I never installed "ubuntu"
 - Run it again - notice how quickly it ran



ssh-ing into a container - fake it...

```
$ docker run -ti ubuntu bash
root@62deec4411da:/# pwd
/
root@62deec4411da:/# exit
$
```

- Now the process is "bash" instead of "echo"
- But its still just a process
- Look around, mess around, its totally isolated
 - rm /etc/passwd – no worries!
 - MAKE SURE YOU'RE IN A CONTAINER!



A look under the covers

```
$ docker run ubuntu ps -ef
```

UID	PID	PPID	C	STIME	TTY
root	1	0	0	14:33	?


TIME	CMD
00:00:00	ps -ef

- Things to notice with these examples
 - Each container only sees its own process(es)
 - Running as "root"
 - Running as PID 1



Docker Images

- Tar file containing a container's filesystem + metadata
- For sharing and redistribution
 - Global/public registry for sharing: DockerHub
- Similar, in concept, to a VM image

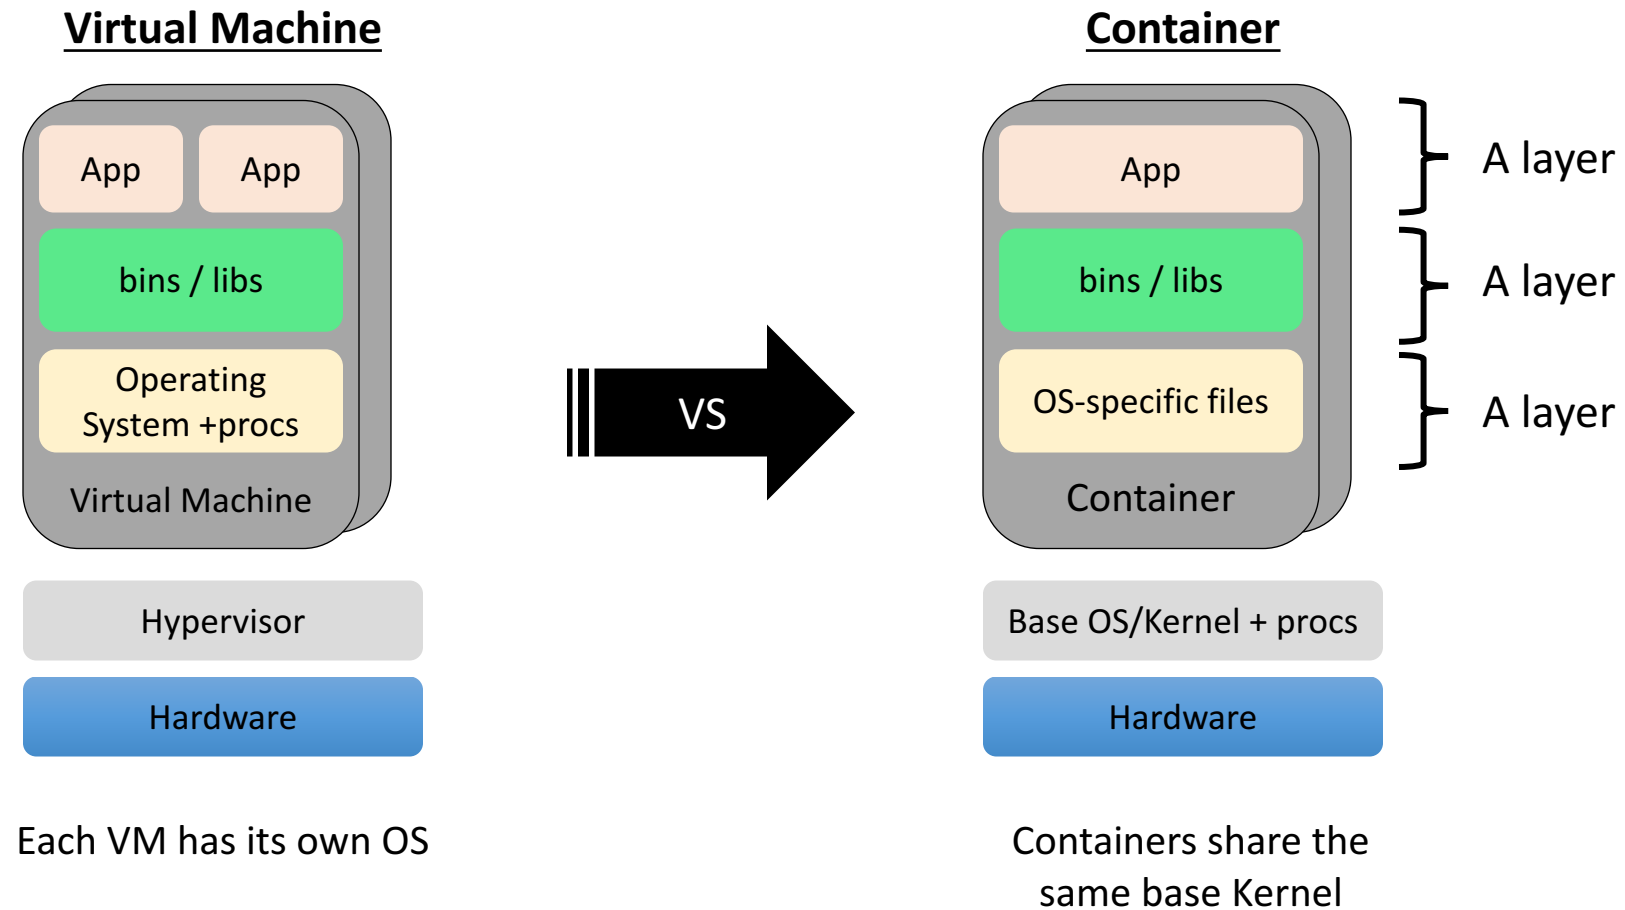


Docker special sauce: Layers

- But first, let's compare VMs and Containers one more time...

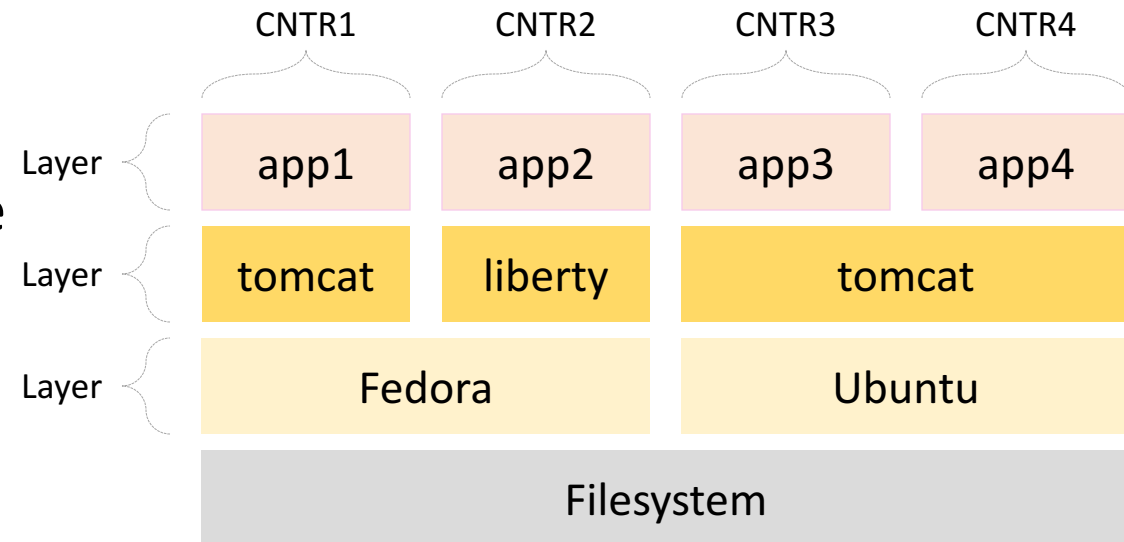


VM vs Container: Notice the layers!



Shared / Layered / Union Filesystems

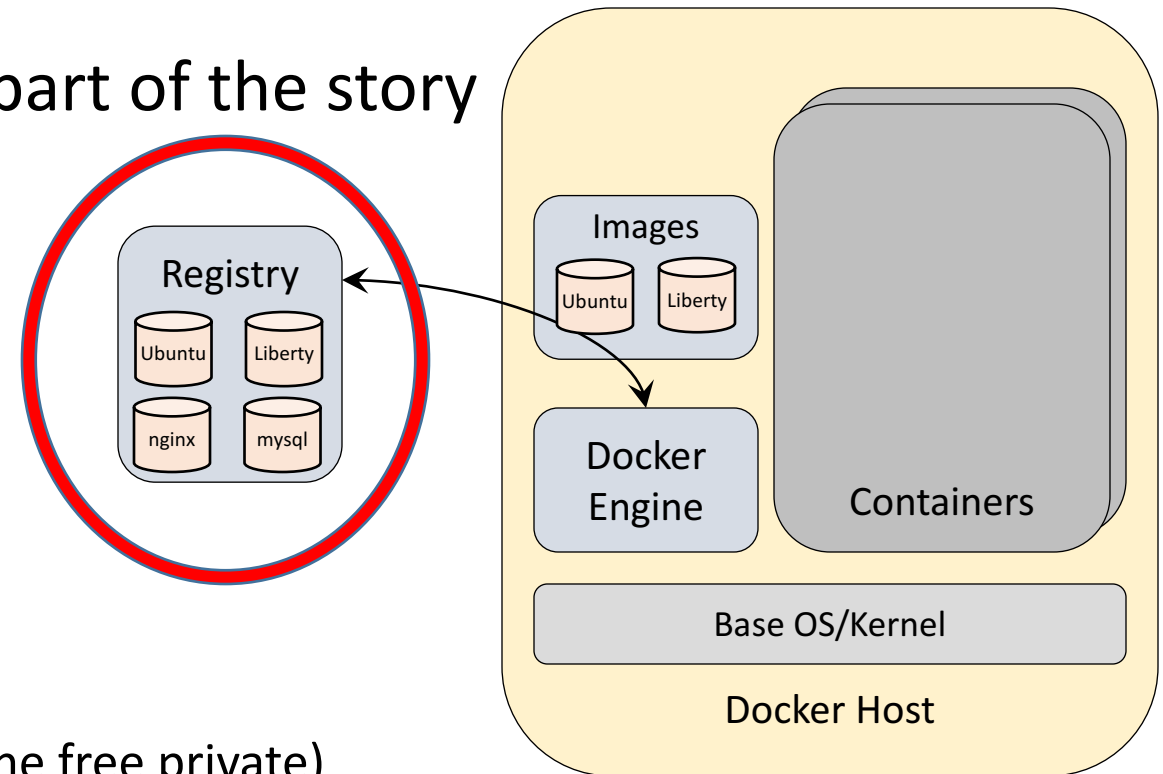
- Docker uses a copy-on-write (union) filesystem
- New files(& edits) are only visible to current/above layers
- Layers allow for reuse
 - More containers per host
 - Faster start-up/download time
- Images
 - Tarball of layers
- Think: Transparencies on projector



Docker Registry

- Creating and reusing images is only part of the story
- Sharing them is the other

- DockerHub - <http://hub.docker.com>
 - Public registry of Docker Images
 - Hosted by Docker Inc.
 - Free for public images, pay for private ones (one free private)
 - By default docker engines will look in DockerHub for images
 - Browser interface for searching, descriptions of images



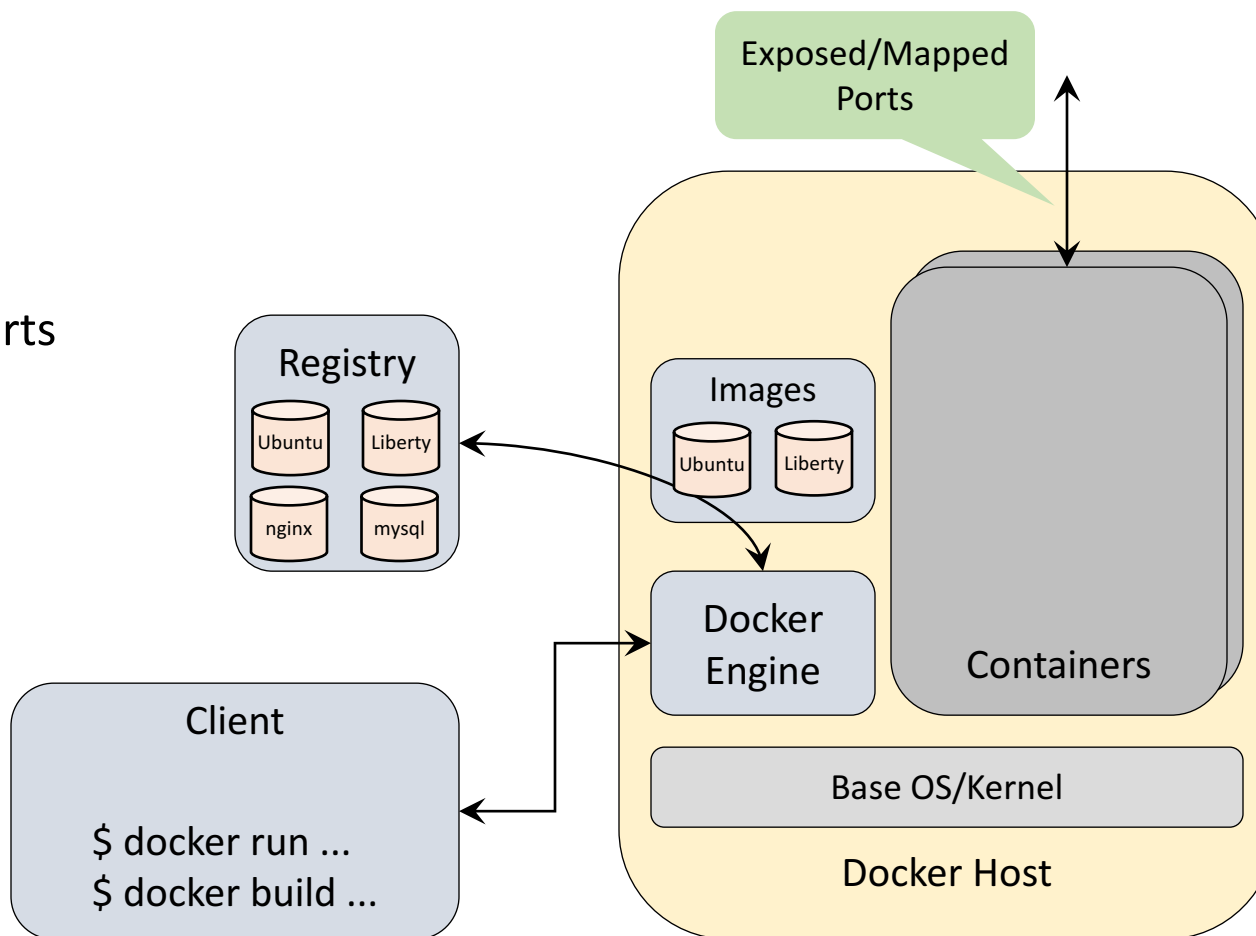
Multi-Architecture Support

- Before: **Docker runs everywhere!** (as long as its x86/Linux)
- Now: Docker daemon has multi-architecture support
 - Docker builds for Power, Z, ARM - Linux
 - Windows CLI built in community, Windows daemon built by Microsoft
- Registry Multi-architecture support is available
 - Engine and OSS Registry code and DockerHub supports it
 - Docker CLI doesn't provide a nice UX yet, but there are tools available
- Engine when pulling down an image:
 - Sends host's arch & OS along with the image tag
 - Registry will find image+arch+OS



Docker Component Overview

- Docker Engine
 - Manages containers on a host
 - Accepts requests from clients
 - REST API
 - Maps container ports to host ports
 - E.g. 80 → 3582
- Images
- Docker Client
 - Drives engine
 - Drives "builder" of Images
- Docker Registry
 - Image DB



Summary

- Docker is just a tool to manage containers
 - Key concepts: Containers, Docker Engine, Images, Registries
- Docker value-add:
 - An excellent User Experience
 - Image Layers
 - Easily shared images via DockerHub
- Why? When compared to VMs:
 - Better resource utilization - CPU, Memory, Disk
 - Faster start-up times
 - Easier tooling/scripting via Docker CLI
- Discussion / Questions?



Quiz!

- What's the difference between a container and an image?
- Answer:
 - An image is a tar of a filesystem
 - A container is a filesystem + a set of processing running in isolation



DevOps Enablement: Configuration Management



Topics

- Configuration Management
- Ensuring developers and stages of the CI pipeline have the correct environment can be a challenge
 - Install variants based on machines
 - Wrong version of products installed



Scenario

- A new developer has joined the team
- They already have:
 - Ubuntu VM with Docker installed
 - "git clone" of the source code for the project:
~ /myapp in the provided VM
- We need to get them up and running as quickly as possible
 - Without installing anything else!



The setup

```
$ cd myapp
```

```
$ cat Makefile
```

```
myapp: myapp.go
```

```
    go build -tags netgo -installsuffix netgo -o myapp myapp.go
```

- There's nothing here about Docker, just a normal compile step



Verify we're missing our dev environment

```
$ make
```

```
go build -tags netgo -installsuffix netgo -o myapp myapp.go
```

```
make: go: Command not found
```

```
Makefile:2: recipe for target 'myapp' failed
```

```
make: *** [myapp] Error 127
```



Solution

- Our IT department has provided a Docker image called "golang"
- This image has the go compiler installed
- Let's use this image to do our build



Using the "golang" image

- Abstractly:
 - Create a new container using the "golang" image
 - Make our source code available inside of the container
 - Build our application in the container
 - Make the executable available outside of the container
 - Otherwise the results will be lost when the container is deleted



Using the "golang" image

- Technically the IT department would setup the Makefile like this **(do not type this)**:

```
docker run golang$(PWD):/src -w /src  
go build -tags netgo -installsuffix netgo -o myapp myapp.go
```

- Summary:
 - **docker run golang** # Creates a container based on "golang"
 - **-v \$(PWD):/src** # Mounts current directory into container at /src
 - **-w /src** # Docker will "cd" to /src before starting process
 - Notice that we didn't modify the normal developer's process, we just wrapped it with Docker
- Quiz: where does the output go?



In action

```
$ make -f Makefile1  
docker run -v /home/user/myapp:/src -w /src golang \  
    go build -tags netgo -installsuffix netgo -o myapp myapp.go
```

- Built "myapp" in the container
- But also stored in current directory due to the mount



Test the Build

```
$ ./myapp 8080
```

```
Will show:
```

```
<pre><b>v1.0 Host: docker Date: 2016-09-04  
05:27:42.582058185 -0700 PDT</b>  
127.0.0.1  
192.168.59.147  
172.17.0.1  
172.19.0.1  
172.18.0.1  
172.20.0.1
```

```
Listening on: 0.0.0.0:8080
```

```
$ curl localhost:8080  
ctrl-c
```

```
# Test it from another window  
# To stop it
```



Summary

- Developer prerequisites:
 - VM + Docker
 - Source code: `git clone ...`
- Concepts:
 - Configuration management/headaches are lessened or removed
 - Developer never had to "install" any dev tooling
 - IT department provided a "standard development environment"
- Discussion / Questions?



Discussion Point

```
$ ls -l
total 5584
-rw-r--r-- 1 user user          44 Sep  3 11:27 Makefile
-rwxr-xr-x 1 root root 5706234 Sep  4 06:51 myapp
-rw-r--r-- 1 user user          887 Sep  3 11:39 myapp.go
```

- "root" owns "myapp"
- This is because we're running as "root" in the container
- Solution: Modify the Makefile to include the user to run as:
docker run **-u 1000:1000** -v \$(PWD):/src -w /src golang \
go build -tags netgo -installsuffix netgo -o myapp myapp.go
 - **\$ make -f Makefile2** if you want to try it, but "**rm -f myapp**" first



Discussion Point

- Did I lie about the prerequisites being just VM+Docker?
- What about "git" ?

```
$ which git # should be empty
$ command -V git # to show the 'alias'
alias git='docker run -ti -u 1000:1000 \
-v /home/user:/tmphome -v $PWD:/wd \
-w /wd -e HOME=/tmphome git'
```

- Git is being run from a container!
- Overrides \$HOME so .gitconfig maps to ~/.gitconfig on host
- Just need to provide the dev with alias - so just a small white-lie
- Containers/isolation/sharing of images isn't just for "real work/services"



DevOps Enablement: Building Value



Topics

- Becoming a creator, and exporter of content, via Docker Images
- Adding value to existing Images
- Sharing this content via Docker Registries
- Becoming part of the value-add chain



Scenario

- Sharing the result of a build with the rest of the CI/CD pipeline
- We have the output of a product build ("myapp" executable)
- We need to build a Docker Image and share it
- NOTE: make sure you're in the "myapp" directory and have build "myapp" executable:

```
$ rm -f myapp  
$ make -f Makefile2
```



Creating a Docker Image - Manually

- Create a Docker Image by "snapshotting" a container
- First we need to create a new container for our application

```
$ docker create ubuntu  
5ed983843bbaef1062096e456e6fd931e6f24e9399d7c801adc7f
```
- Now let's copy our executable into it:

```
$ docker cp myapp 5ed98:/myapp
```
- Finally, snapshot the container as a Docker Image - called "myapp"

```
$ docker commit -c "entrypoint /myapp" 5ed98 myapp  
sha256:7c640789dae5607c868a56883189d6c72478eff1080a67
```



Test the Image

```
$ docker run -ti myapp
```

Will show:

```
<pre><b>v1.0 Host: 165dcbc3e6f8   Date: 2016-09-05 02:47:50.2...</b>
127.0.0.1
172.17.0.2
```

Listening on: 0.0.0.0:80

- In another window:

```
$ curl 172.17.0.2
```

```
<pre><b>v1.0 Host: b8d73b85cc04   Date: 2016-09-05 02:53:49.803922...</b>
127.0.0.1
172.17.0.2
```

- Stop the app by pressing: **ctrl-c** in first window



Discussion

- Can we expose this container at the host level so others can access it?
- Yes, by mapping port 80 in the container to a unique port on the host

```
$ docker run -d -ti -p 9999:80 myapp  
4b08d035deb6135eff60babd1368ab47c0c1f1d09a8ddf3f9417e7e4c4
```

```
$ curl localhost:9999  
<pre><b>v1.0 Host: 4b08d035deb6   Date: 2016-09-05  
03:14:31.713...</b>  
127.0.0.1  
172.17.0.2
```

```
$ docker rm 4b08  
Failed to remove container ...
```

```
$ docker rm -f 4b08  
4b08
```



Creating a Docker Image - With Docker Build

- Docker provides a "build" feature
- Uses a "Dockerfile"
 - Like a "Makefile", a list of instructions for how to construct the container

```
$ cat Dockerfile  
FROM ubuntu  
ADD myapp /  
EXPOSE 80  
ENTRYPOINT /myapp
```



Creating a Docker Image - With Docker Build

```
$ docker build -t myapp .  
Sending build context to Docker daemon 5.767 MB  
Step 1/4 : FROM ubuntu  
----> ff6011336327  
Step 2/4 : ADD myapp /  
----> b867e19a859b  
Removing intermediate container ea699ecc51a0  
Step 3/4 : EXPOSE 80  
----> Running in 85c240f03ae9  
----> 5d8e53bbf9e4  
Removing intermediate container 85c240f03ae9  
Step 4/4 : ENTRYPOINT /myapp  
----> Running in f318d82c2c38  
----> 684c6c2572ff  
Removing intermediate container f318d82c2c38  
Successfully built 684c6c2572ff
```



Test the image - With Auto-Port Allocation

```
$ docker run -tidP myapp
```

```
469221295fae1b57615286ec7268272e3d3583c12ea66e14b2
```

```
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
469221295fae	myapp	"/bin/sh -c /myapp"	4 seconds ago	Up 4
seconds	0.0.0.0:32768->80/tcp	clever_ardinghelli		

```
$ curl localhost:32768
```

```
<pre><b>v1.0 Host: 469221295fae Date: 2016-09-05 03:34:49....</b>  
127.0.0.1  
172.17.0.3
```

```
$ docker rm -f 469
```

```
469
```



Sharing the Image

- The "myapp" image is only in our local image cache
- To distribute it we need to upload it to a shared registry
- First, let's start a registry - locally, but pretend its public
\$ regStart



Naming Images

- Before uploading an image, its name must include the registry
- General syntax of image names:
 - `[[registry/] [namespace/]] name [:tag]`
 - E.g. `docker:5000/myapp:1.0` # "docker" is our hostname
- Registry: host:port - presence of ":" disambiguates from "namespace"
- Namespace: user, owner
- Tag: typically a version string - defaults to "latest"



Preparing our Image

```
$ docker build -t docker:5000/myapp:1.0 .  
Sending build context to Docker daemon 5.767 MB  
Step 1/4 : FROM ubuntu  
----> ff6011336327  
Step 2/4 : ADD myapp /  
----> Using cache  
----> b867e19a859b  
Step 3/4 : EXPOSE 80  
----> Using cache  
----> 5d8e53bbf9e4  
Step 4/4 : ENTRYPOINT /myapp  
----> Using cache  
----> 684c6c2572ff  
Successfully built 684c6c2572ff
```

- Alternative:
\$ docker tag myapp docker:5000/myapp:1.0



Pushing the Image

```
$ docker push docker:5000/myapp:1.0
```

The push refers to a repository [docker:5000/myapp]

5d1c38831713: Pushed

447f88c8358f: Pushed

df9a135a6949: Pushed

dbaa8ea1faf9: Pushed

8a14f84e5837: Pushed

latest: digest: sha256:71f76c1b360e340614a52bcfef2cb78d8f0aa3604 size: 1363

- Image is in the registry and can be used by other parts of the pipeline

```
$ docker run -ti docker:5000/myapp:1.0
```

```
$ docker pull docker:5000/myapp:1.0
```



Summary

- Can build an image via "docker build" or manually/scripted
- Share image via registries with the rest of the CI/CD pipeline
 - Images are not just for "products"
 - Move your testcases through the CI/CD pipeline as well
 - Anything you want to share can be in an Image
- Discussion / Questions?



Discussion Point

- Our Dockerfile started with: **FROM ubuntu** do we really need Ubuntu ?
- No, there is nothing in our app that uses the operating system

```
$ docker images | grep myapp
```

```
docker:5000/myapp 1.0 684c6c2572ff 7 hours ago 193.7 MB
```

- Instead our Dockerfile could use: **FROM scratch**
 - Let's do that so our image is smaller

```
$ docker build -f Dockerfile2 -t docker:5000/myapp:1.0 .
```

```
$ docker images | grep myapp
```

```
docker:5000/myapp 1.0 9b604f2e42da 7 seconds ago 5.84 MB
```

```
$ docker push docker:5000/myapp:1.0 # update our registry
```



Discussion Point

- What are **some** of the other instructions can we have in a Dockerfile?
 - RUN
 - HEALTHCHECK
 - COPY/ADD
 - CMD & ENTRYPOINT
 - LABEL
 - ENV/ARG
 - VOLUME
 - USER
 - WORKDIR



DevOps Enablement: Scaling and Reliability



Topics

- Take a single container/image and spread it
 - To provide redundancy, reliability and performance
- Can be complex to manage production-level deployments
 - Clustering
 - High-availability
 - Scaling
 - Versioning - with no downtime



Scenario

- Given an application/image we need to deploy it
 - With high-availability
- Then we need to upgrade it
 - With zero down-time



Docker Orchestration

- In v1.12 Docker introduced the notion of a "Service"
- A scaled cluster of Docker containers based on the same Image
- Across a cluster of Docker Hosts
- Built into the Docker engine

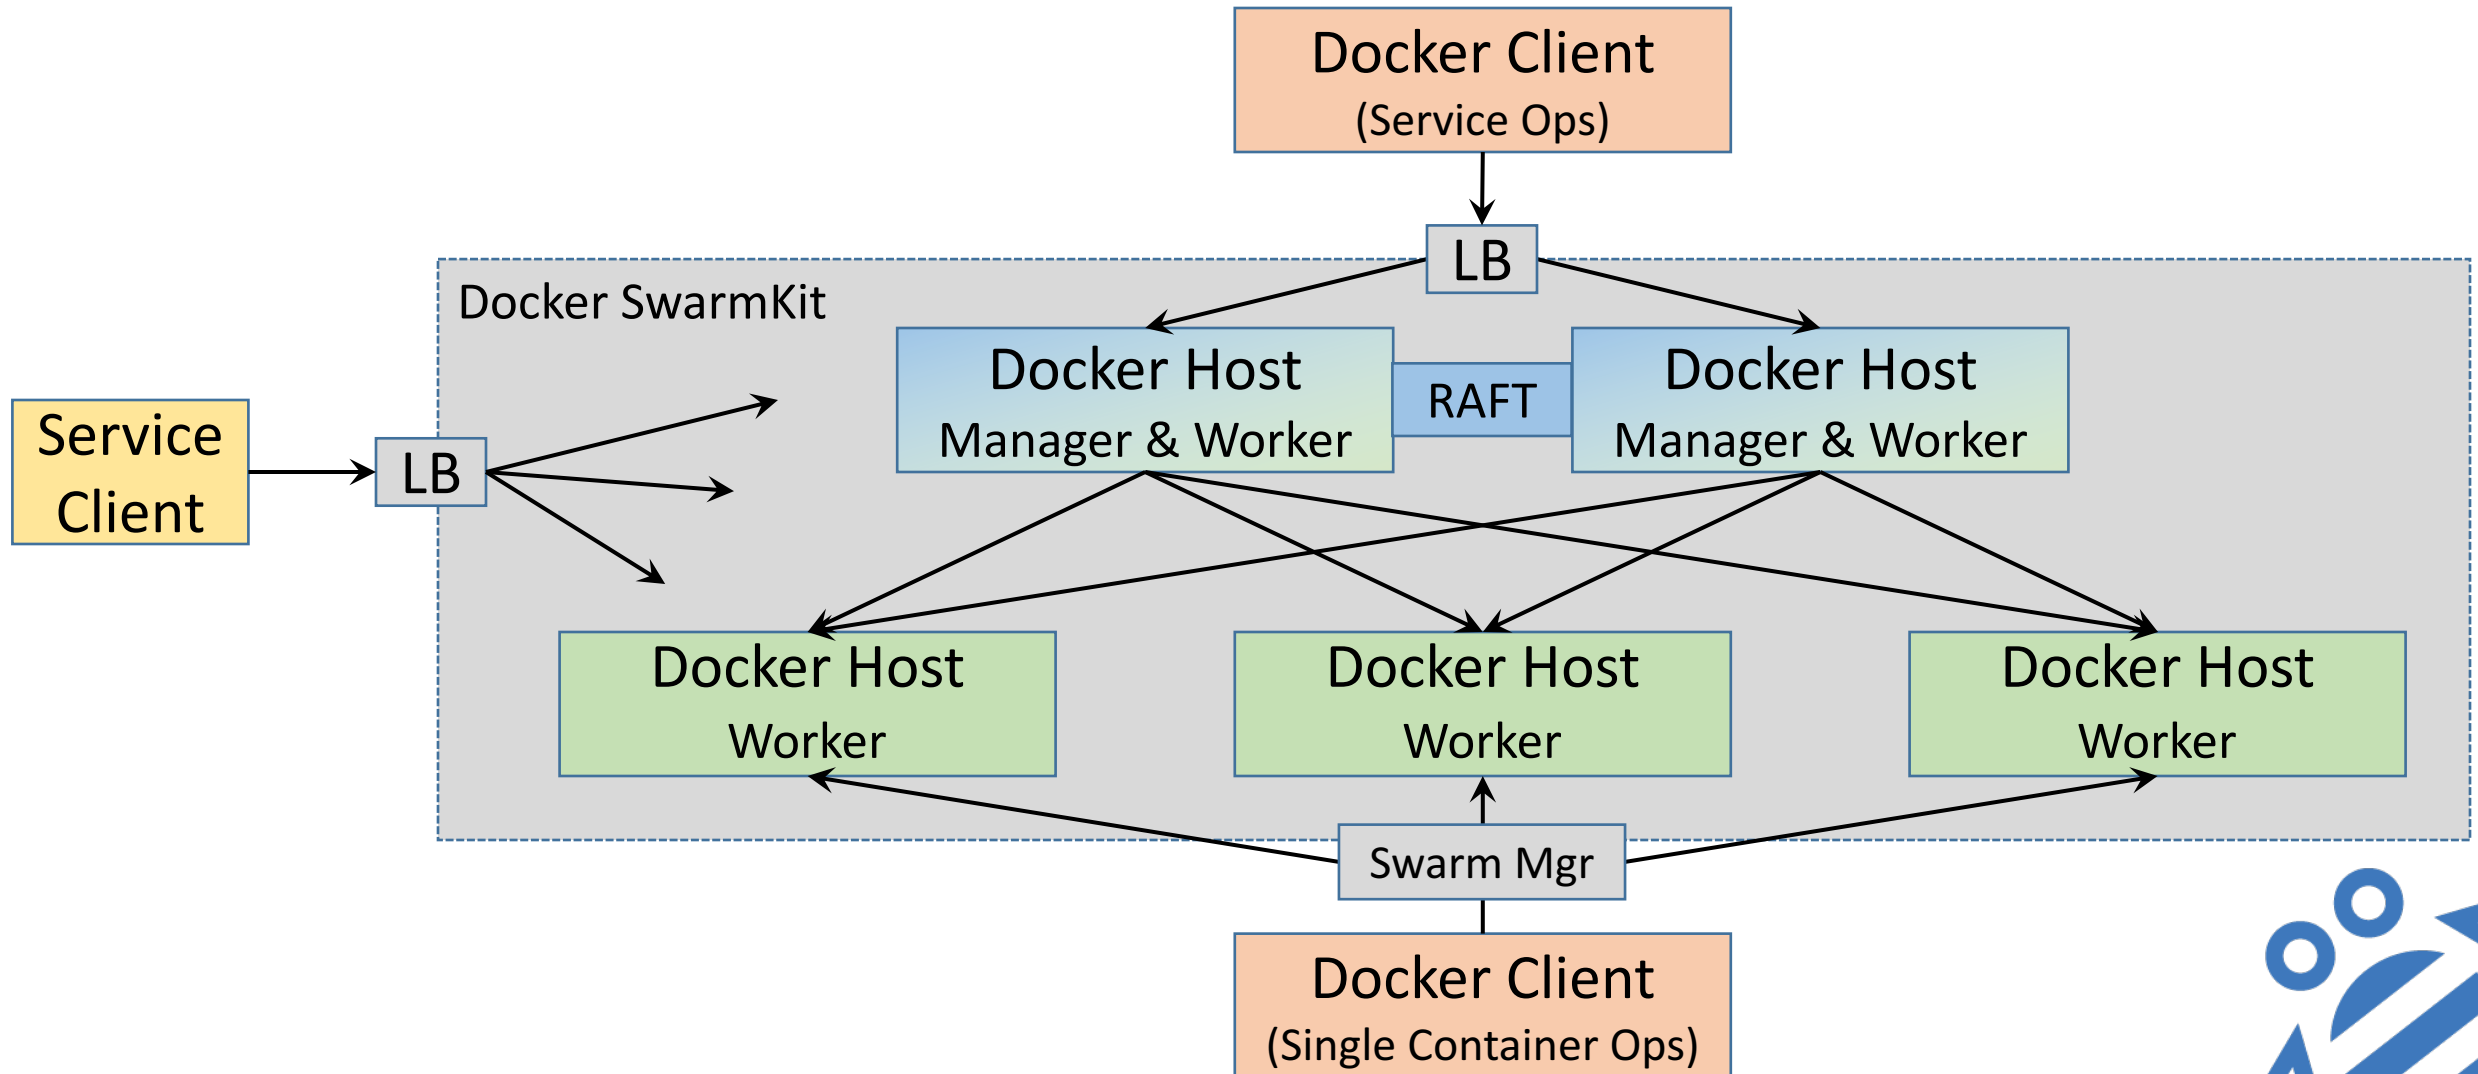


Terminology

- Worker
 - A Docker Host capable of accepting requests to manage single containers
 - New name for the existing Docker host/engine
- Manager
 - A Docker Host capable of accepting requests to manage "services"
 - Also acts as a Worker
- As a Docker client you need to know which to talk to
 - Single container management vs Service management



Workers, Managers and the "Mesh"



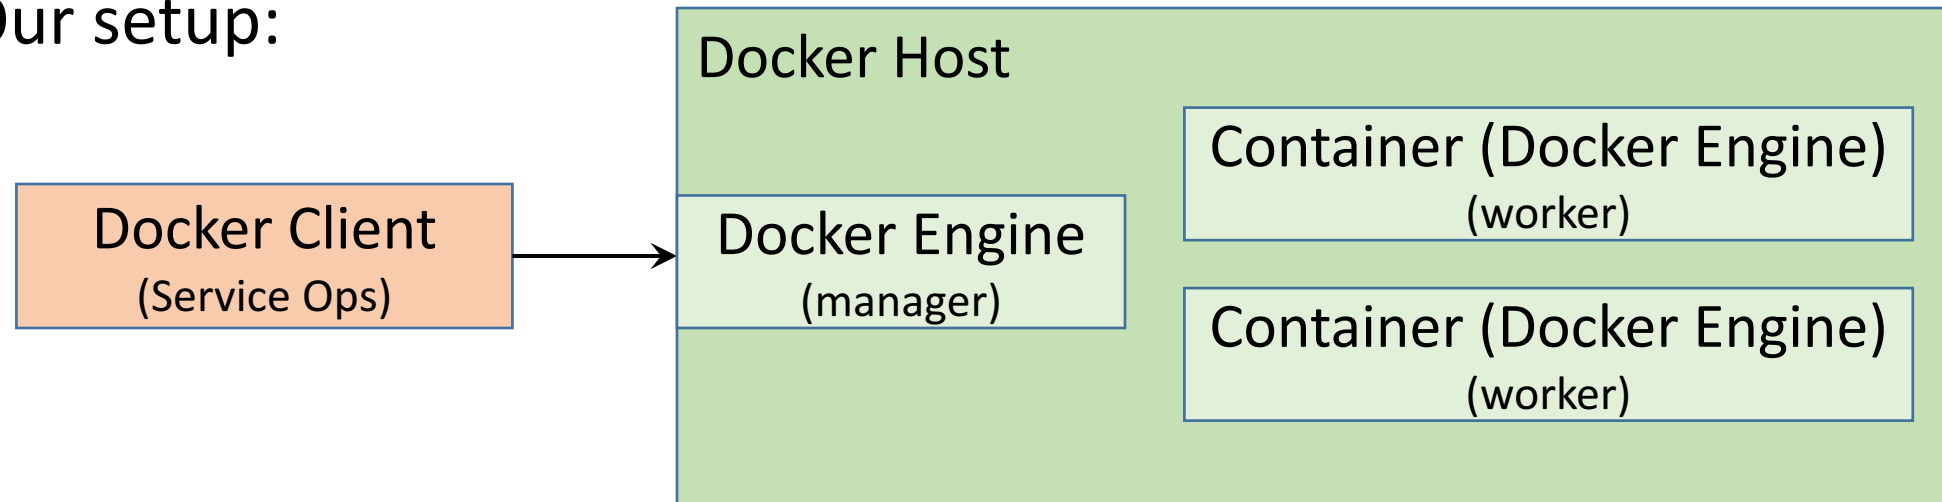
Extra Details

- Each "service" gets its own Virtual IP
- DNS resolution will map "service name" to "VIP"
- Introduces a declarative model with eventual consistency
- RAFT : in-memory distributed state store
- Shared port space across **all** nodes



Part 1 - Deploy our Docker Cluster

- Setup a Docker SwarmKit Cluster to host our application
- Our setup:



\$ swarmDemo1



Part 2 - Deploy v1.0 of our app

- Deploys the application to the swarm
- Scale it to 3 instances

\$ swarmDemo2



Part 3 - Build a v2.0 of the app

```
$ docker build -f Dockerfile3 -t docker:5000/myapp:2.0 .
```

- Dockerfile3:

```
FROM scratch  
ADD myapp /  
EXPOSE 80  
ENV APP_VER=2.0  
ENTRYPOINT /myapp
```

```
$ docker push docker:5000/myapp:2.0
```



Part 4 - Upgrade the app

In one window run:

```
$ docker exec -ti node3 watch -d -n 1 curl -s 127.0.0.1
```

This will continually hit the app, but do so by sending the request to node3. Notice the version string starts out as "1.0".

In another window update the service:

```
$ docker service update --image=docker:5000/myapp:2.0 myapp
```

Notice the version string will eventually change to "2.0"

```
$ swarmDemo9          # to clean up
```



Summary

- Deploying, scaling and upgrading an app can be complicated
- Docker manages this for us with an "eventual consistency" model
 - We provide the "desired state"
- Discussion / Questions?

